

Dummy PDF Content

This is a small document used for the RAG component of the AI Agent Pipeline assignment.

Section 1: Vector Databases

A **vector database** is a database optimized for storing and querying high-dimensional vector data. These vectors, often called embeddings, are numerical representations of text, images, or audio. They are crucial for semantic search and Retrieval-Augmented Generation (RAG).

Qdrant is an open-source vector similarity search engine. It is designed to be highly scalable and provides a feature-rich platform for building advanced vector search applications. Qdrant is especially good for managing large-scale vector collections with metadata filtering.

Section 2: Agentic Workflow

An **agentic workflow** uses an LLM not just to generate text, but to make decisions about which steps to take next. LangGraph is built for this, allowing the creation of cyclic graphs where the agent can loop or choose between tools.